

# Next Generation of Platform Engineering Using Kcp and Crossplane

*Lovro Sviben & Simon Bein*



# Simon Bein

Senior Software Engineer @ Kubermatic

[github.com/SimonTheLeg](https://github.com/SimonTheLeg)

[linkedin.com/in/simon-bein](https://linkedin.com/in/simon-bein)

[twitch.tv/simonbuilds](https://twitch.tv/simonbuilds)

I also maintain [konf](#), the “best kubeconfig manager out there” (citation needed)



# Lovro Sviben

Senior Distributed Systems Engineer @  
Upbound

[github.com/lsviben](https://github.com/lsviben)

[linkedin.com/in/lovro-sviben](https://linkedin.com/in/lovro-sviben)



# Common Understanding of Platform Engineering



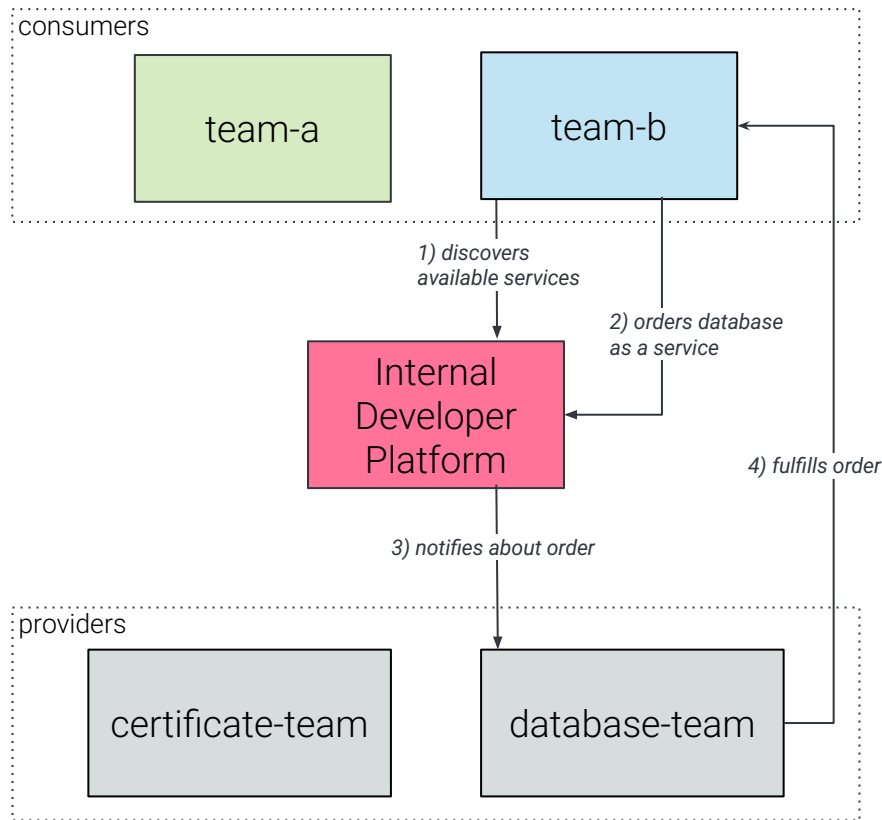
# **Within this talk, Platform Engineering is defined as**

Creation and Maintenance of an automated,  
self-service platform for services inside a company



# Enough Theory - Here's an Example

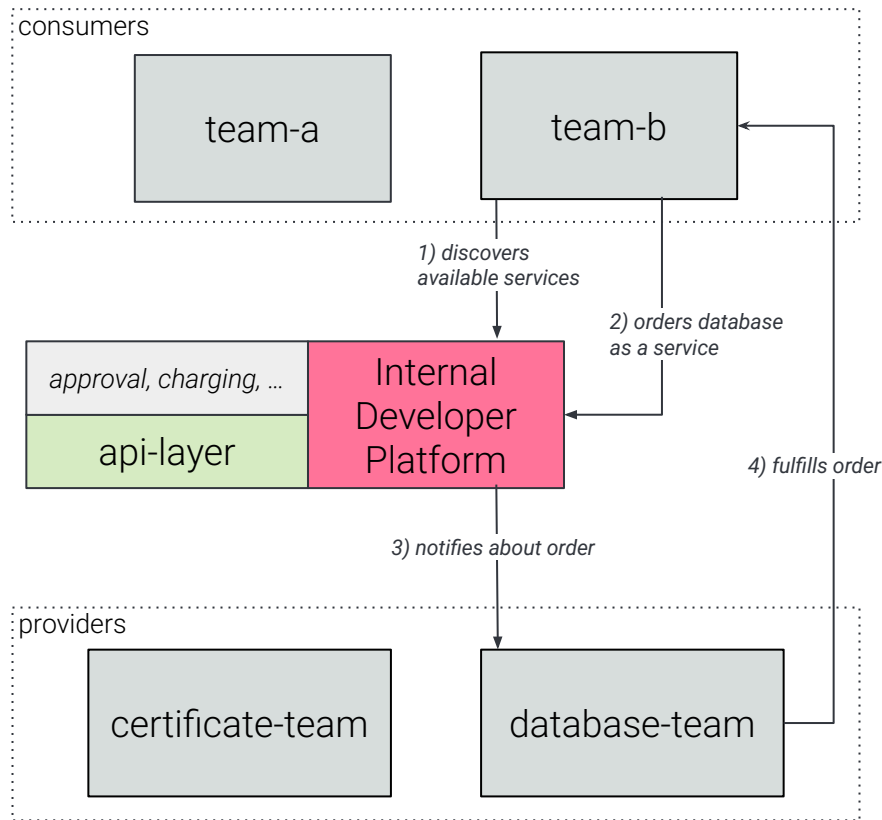
Creation and Maintenance of an automated, self-service platform for services inside a company



# This Talk Focuses on the API Layer of such a Platform

API Layer is responsible for:

- Publishing and discovery of services
- Lifecycle and ownership of Service APIs
- Standardized Data interface for other components



# Why Kubernetes is interesting as an Api-Layer





# (1) Sophisticated API Conventions within Kubernetes

/apis/<group>/<version>/[namespaces/<namespace>/]<resourcetype>[/<name>]

APIs in Kubernetes  
are **grouped**.

Each API group is  
also **versioned**.

Resources are optionally  
**namespaced**.

Resources have a specific  
**resource type** that defines  
their schema.

Resources are  
uniquely **named**.



## (2) The Kubernetes API is extendable

We can extend APIs available in  
the kube-apiserver using CRDs

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: shirts.stable.example.com
spec:
  group: stable.example.com
  scope: Namespaced
  names:
    kind: Shirt
  versions:
    - name: v1
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                color:
                  type: string
                size:
                  type: string
                ...
```



## Some Examples



# The Kubernetes API is pretty awesome!

(that's it. That's the ~~tweet post~~ slide)



## But ...

- APIs (CRDs) are cluster-scoped, so everyone shares them



# Let's Give Everyone a Cluster!



## But ...

- APIs (CRDs) are cluster-scoped, so everyone shares them
- Starting a new cluster with its own api-server is time and resource intensive



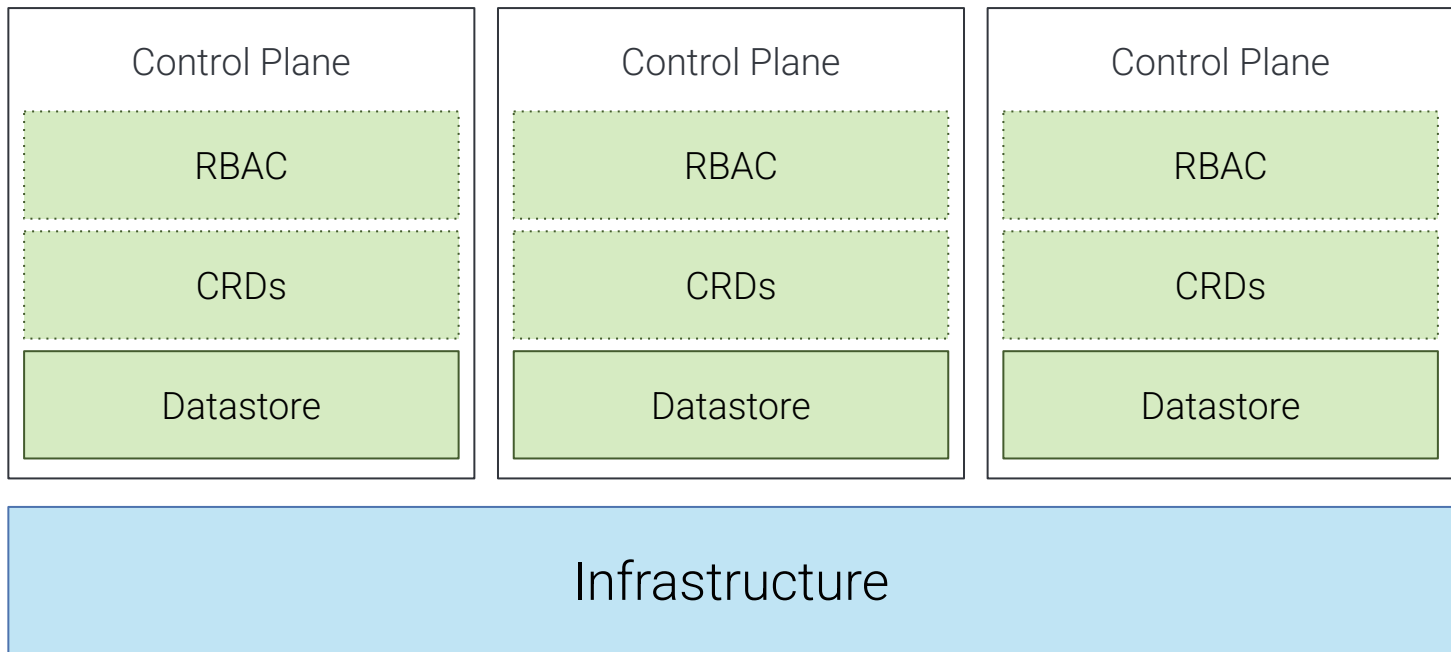
# Lightweight Clusters

to the rescue?





## Hosted Control Planes



## But ...

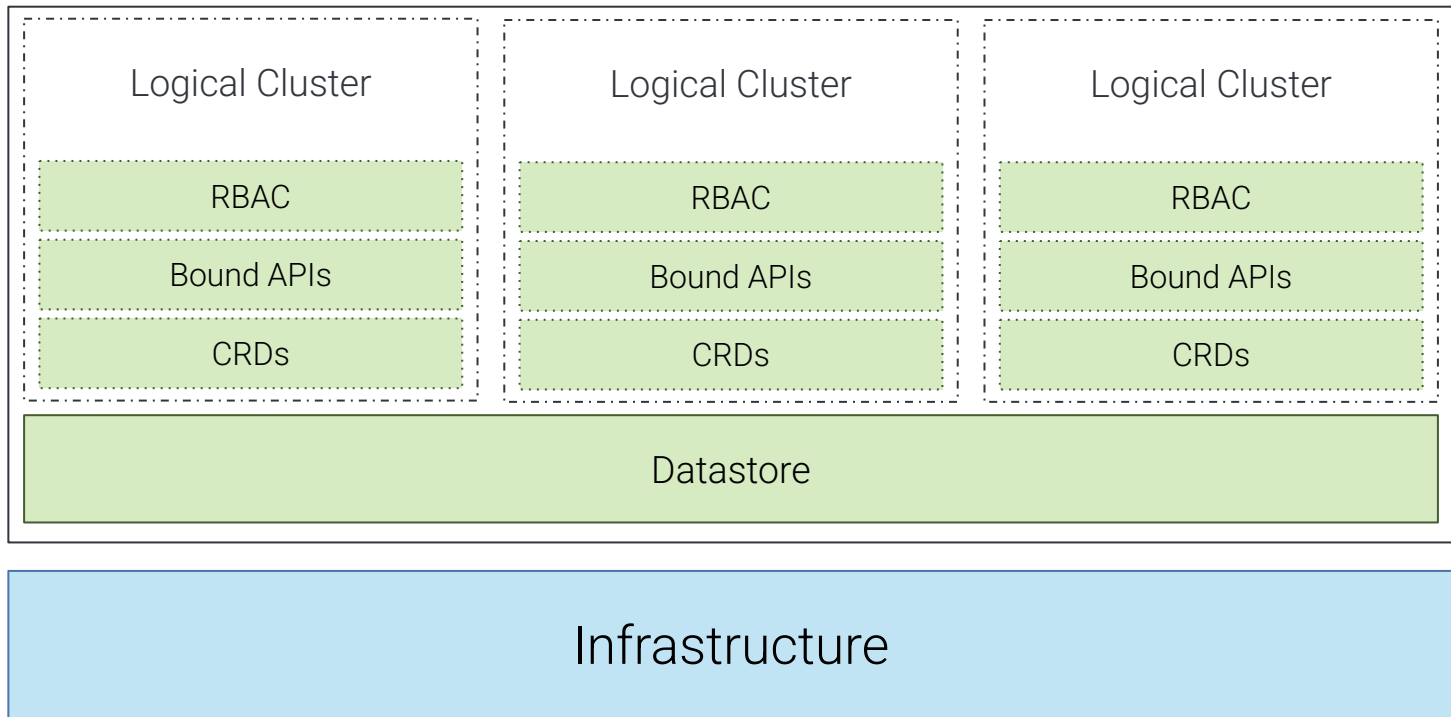
- APIs (CRDs) are cluster-scoped, so everyone shares them
- Starting a new cluster with its own api-server is time and resource intensive
- Sharing apis between a large number of clusters is cumbersome



# What if Control Planes share data?



# “Logical” Clusters



# **This brings us to kcp**





**CLOUD NATIVE**  
**COMPUTING FOUNDATION**

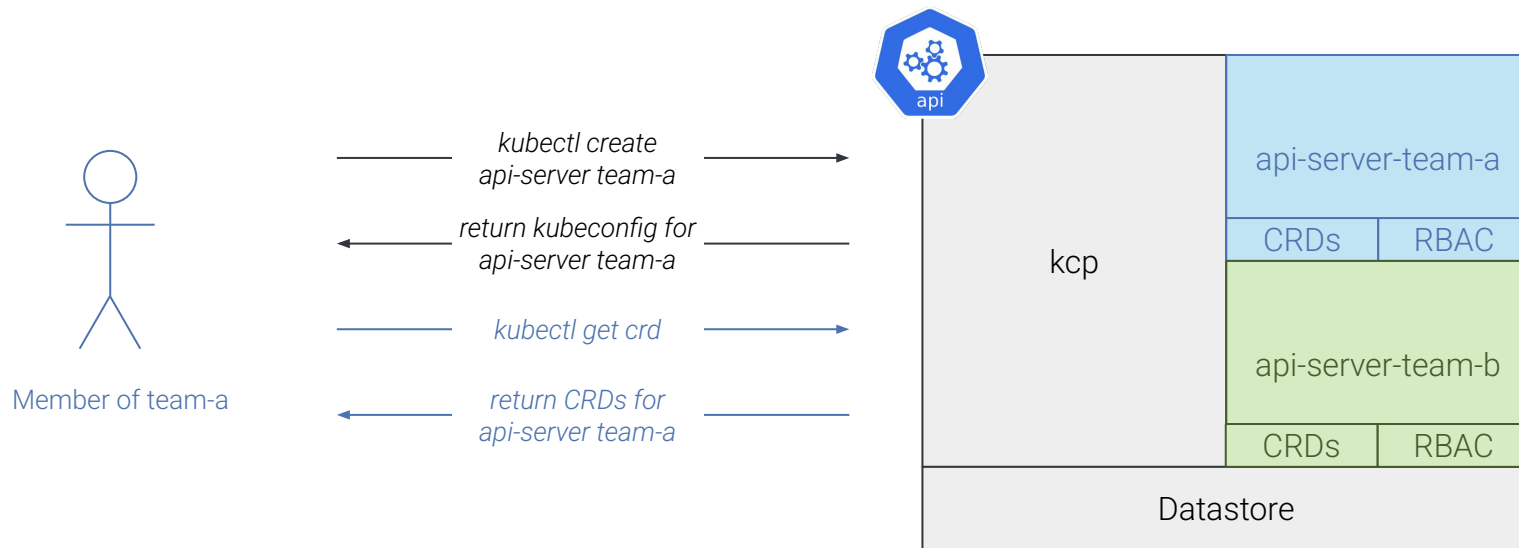
**Sandbox project**  
(since end of 2023)



**“A horizontally scalable  
control-plane for  
Kubernetes-style APIs”**

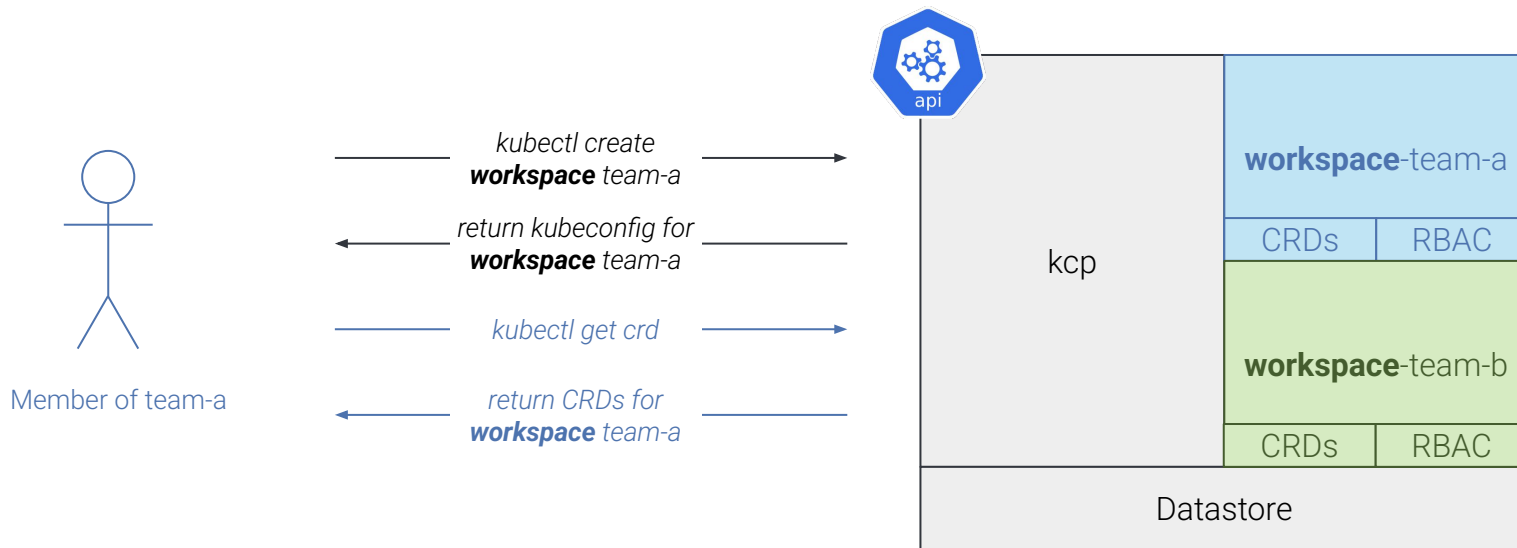


# A practical example of interacting with kcp





# Instead of “api-servers”, we call them “workspaces”



# Workspace

A multi-tenancy **unit of isolation** in kcp.

Each workspaces has its own available **API resource types**.

API **objects** are not shared across workspaces.

Delegation of **administrative permissions** to workspace owners.

Workspaces are **cheap**.



# Workspaces are organized in a tree


`https://kcp:6443/clusters/root`

---

`https://kcp:6443/clusters/root:org-a`

---

`https://kcp:6443/clusters/root:org-a:team-a`



# Practical Example

```
$ kcp start --bind-address=127.0.0.1  
... lots of log output
```

```
$ export KUBECONFIG=.kcp/admin.kubeconfig
```

```
$ kubectl create configmap foo  
configmap/foo created
```

```
$ kubectl create workspace team-a  
Workspace "team-a" (type root:organization) is ready to use.
```

```
$ kubectl get ws
```

NAME	TYPE	REGION	PHASE	URL	AGE
team-a	team		Ready	https://...	3m23s

```
$ kubectl ws :root:team-a  
Current workspace is "root:team-a" (type root:team).
```

```
$ kubectl get configmap foo  
Error from server (NotFound): configmaps "foo" not found
```



# Workspaces allow filesystem-like navigation

```
$ kubectl ws .  
Current workspace is "root".
```

```
$ kubectl get ws
```

NAME	TYPE	REGION	PHASE	URL	AGE
org-a	organization		Ready	https://...	69d
org-b	organization		Ready	https://...	65d

```
$ kubectl ws org-a  
Current workspace is "root:org-a" (type root:organization).
```

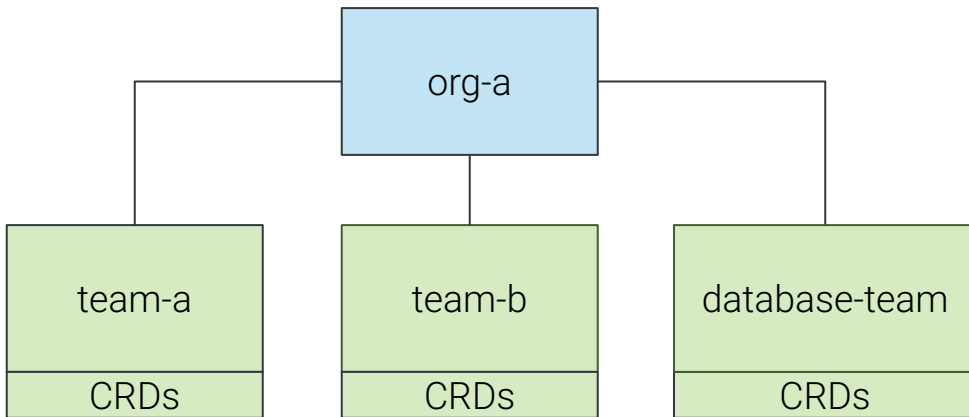
```
$ kubectl get ws
```

NAME	TYPE	REGION	PHASE	URL	AGE
team-a	team		Ready	https://...	3m23s
team-b	team		Ready	https://...	3m18s


```
$ kubectl ws team-a  
Current workspace is "root:org-a:team-a" (type root:team).
```



## When we go back to our example...



# ... the reality is closer to this



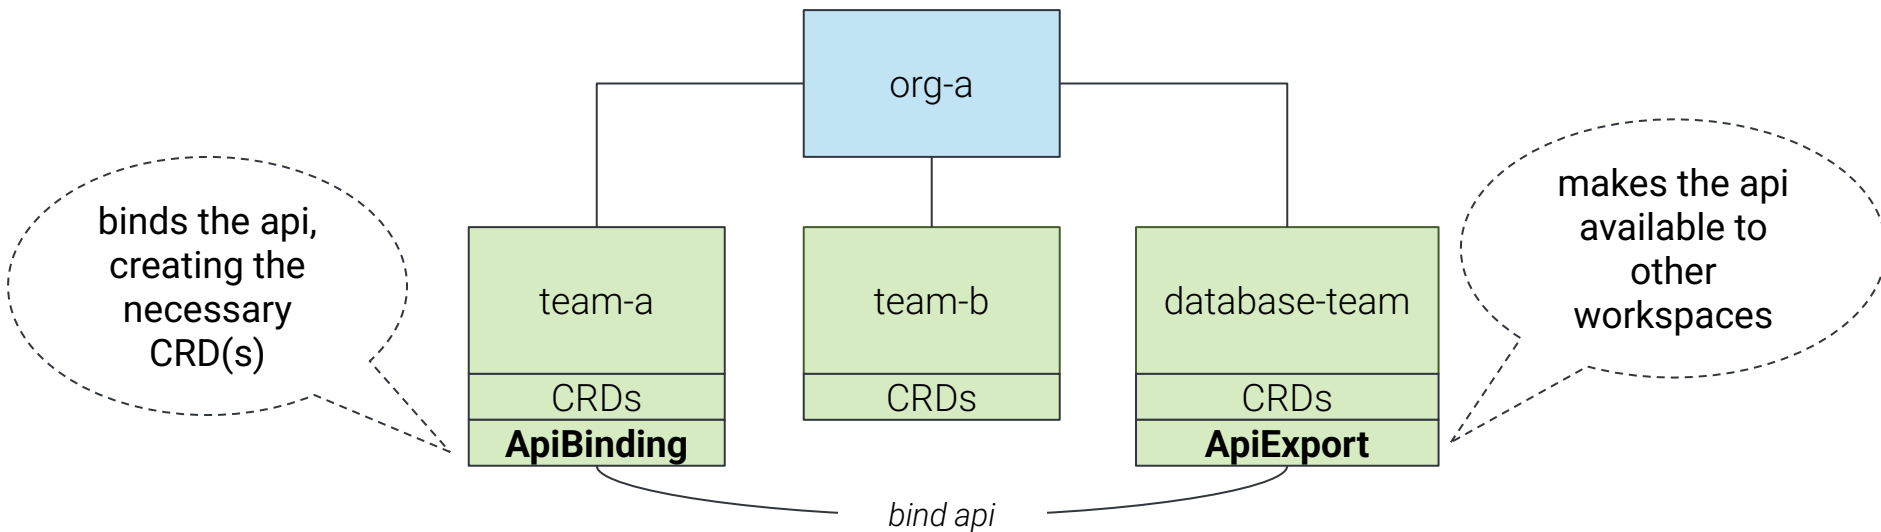
# The API Marketplace

Sharing is Caring





# How this can be achieved in kcp: ApiExport & ApiBinding



Infrastructure + Service teams are not in the business of making APIs discoverable and consumable.

**Platform teams are.**



# Create APIs with APIExports



# APIExport

```
apiVersion: apis.kcp.io/v1alpha1
kind: APIExport
metadata:
  name: databases.demo.example.com
spec:
  latestResourceSchemas:
    - v1.databases.demo.example.com
    - v1.databaseDrivers.demo.example.com
```

Resource schemas define  
resources, just like CRDs.



# Enable APIs with APIBindings



# Powered by APIBindings

```
$ kubectl get apibindings
```

NAME	AGE	READY
<b>tenancy.kcp.io-3wb5h</b>	<b>30d</b>	<b>True</b>
topology.kcp.io-cua3o	30d	True

```
apiVersion: apis.kcp.io/v1alpha1
kind: APIBinding
metadata:
  name: tenancy.kcp.io-3wb5h
spec:
  reference:
    export:
      name: databases.demo.example.com
      path: root:database-team
```

This references an APIExport in a different workspace!

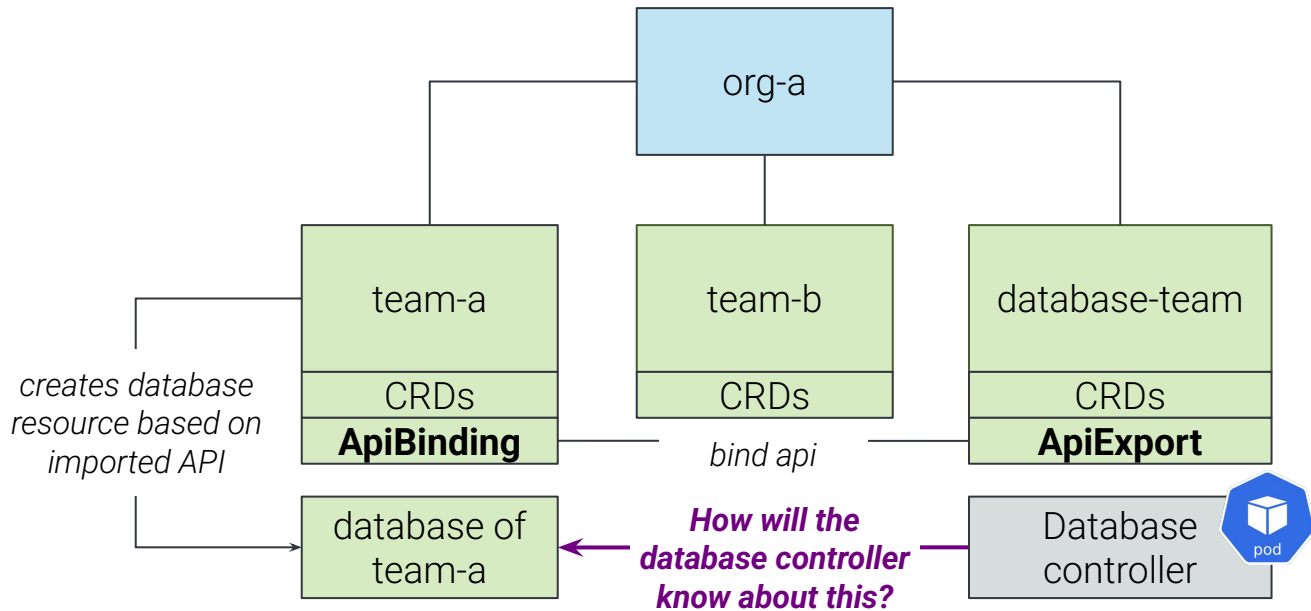


# Discovery of API-Consumers

Keeping Track of Orders



# Actually the reality is even more complicated

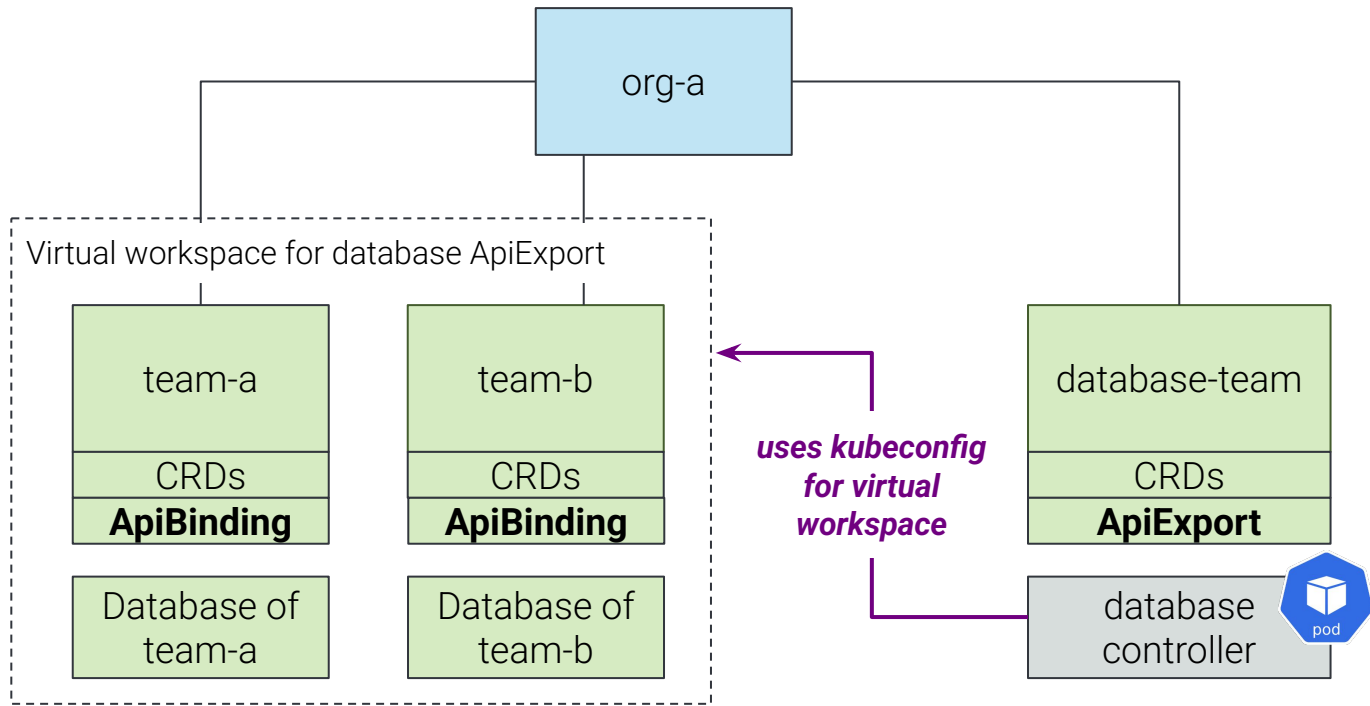




# The APIExport Virtual Workspace to the Rescue



# A virtual workspace provides a computed view of parts of one or multiple workspaces



# APIExport Virtual Workspace - Good to know

- Virtual Workspaces are not a CustomResource Type, you cannot directly create one, but kcp provides them for you
- APIExport Virtual Workspace provides a unique URL for its export, which can be used as the server URL in a k8s controller
- There are more virtual workspaces available in kcp



# Intro to Crossplane



# What is Crossplane

- Your cloud native **control plane**
  - Provision/manage **all** of your resources
- **Compose** those resources into high level **abstractions**
  - Give your developers self-service provisioning
- Kubernetes is a great control plane for containers
  - Crossplane teaches it how to manage **everything** else
- Cloud providers have used control planes for years
  - Now it's your turn to build your own!



# Build your own

- Assemble granular resources. Crossplane has a marketplace with providers managing various resources.
- Expose as higher level self-service API for your app teams
  - **Compose** GKE, NodePool, Network, Subnetwork
  - **Offer** as a simple **Cluster** abstraction (API) with limited config for developers to self-service
- Hide infrastructure complexity and codify a “golden path”
- All with K8s API - compatible with kubectl, GitOps, etc.
- No code **required**



# Composite

```
apiVersion: apiextensions.crossplane.io/v1
kind: CompositeResourceDefinition
metadata:
  name: nosqls.database.example.com
spec:
  group: database.example.com
  names:
    kind: NoSQL
    plural: nosqls
  versions:
  - name: v1alpha1
    served: true
    referenceable: true
    schema:
      openAPIV3Schema:
        type: object
        properties:
```

First create Composite Resource Definition (XRD) to declare our custom platform API

API Group

Standard OpenAPI v3 schema



# Composition

```
apiVersion: apiextensions.crossplane.io/v1
kind: Composition
metadata:
  name: nosqls.database.example.com
spec:
  compositeTypeRef:
    apiVersion: database.example.com/v1alpha1
    kind: NoSQL
  mode: Pipeline
  pipeline:
  - step: generate-resources
    functionRef:
      name: function-acme-func
      input: {}
  - step: filter-resources
    functionRef:
      name: function-filter
      input: {}
```

Then we define a  
Composition that  
defines the XRD

The XRD this  
Composition is for

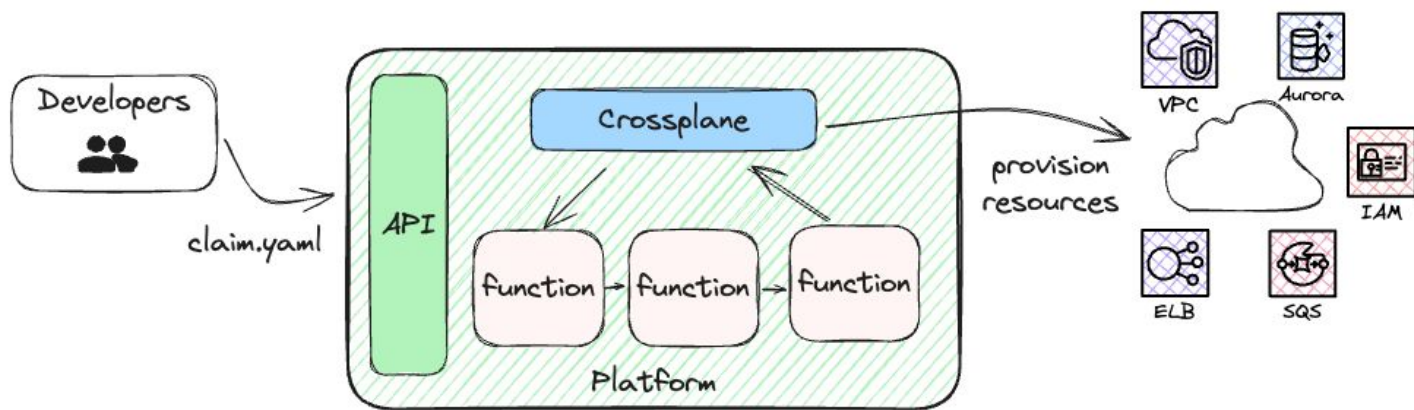
Pipeline of functions to  
execute that will generate  
the managed resources





# How do Functions work

- Run a pipeline of simple functions to compose resources
- Written in your language of choice
- Focus only on your unique logic
- Crossplane does the heavy lifting of resources CRUD, reconciling, finalizers, owner refs, etc.



# Crossplane v2 is here

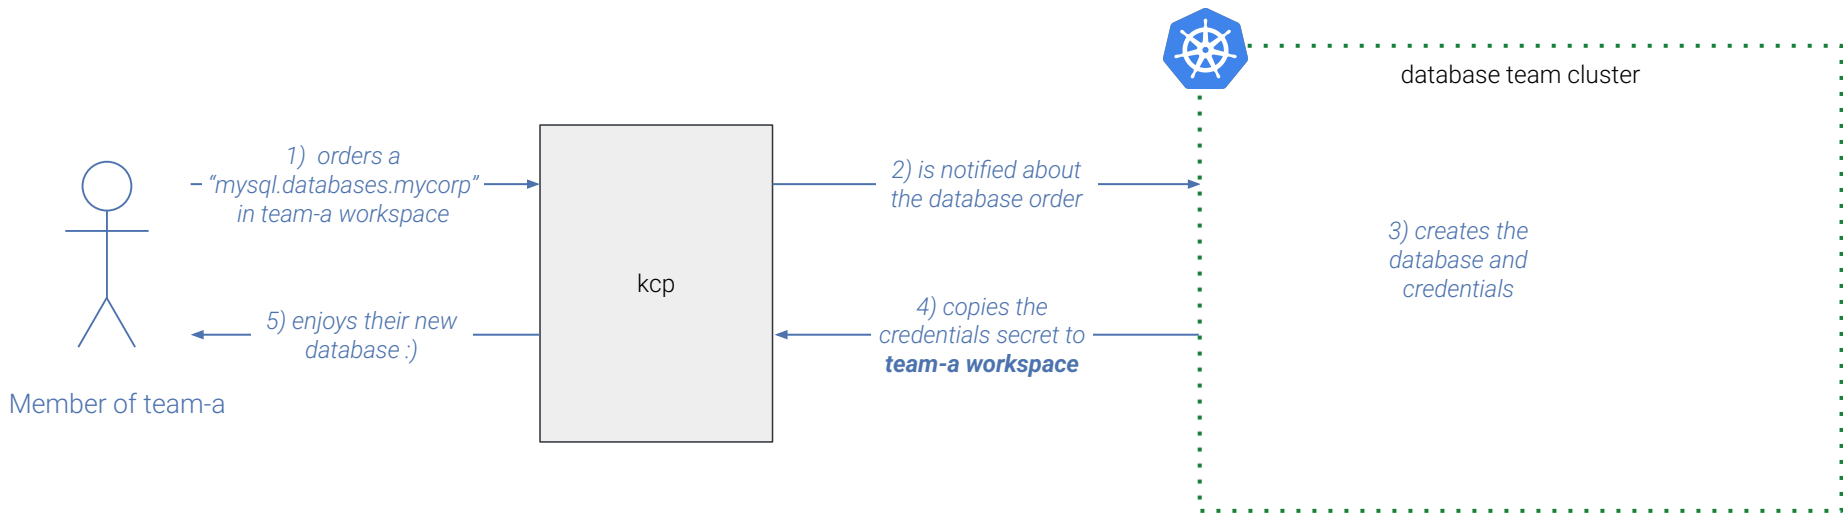
- **Crossplane v2 is more useful, more intuitive, and less opinionated**
- Three major changes:
  - Composite resources can now be namespaced
  - Managed resources are now namespaced
  - Composition supports any Kubernetes resource
- Operations - a new way to run operational tasks
- ManagedResourceDefinitions - a new way to control CRD sprawl
- Crossplane v2 is better suited to building control planes for **applications**, not just infrastructure



# Demo Time



# What do we want to achieve? - High Level



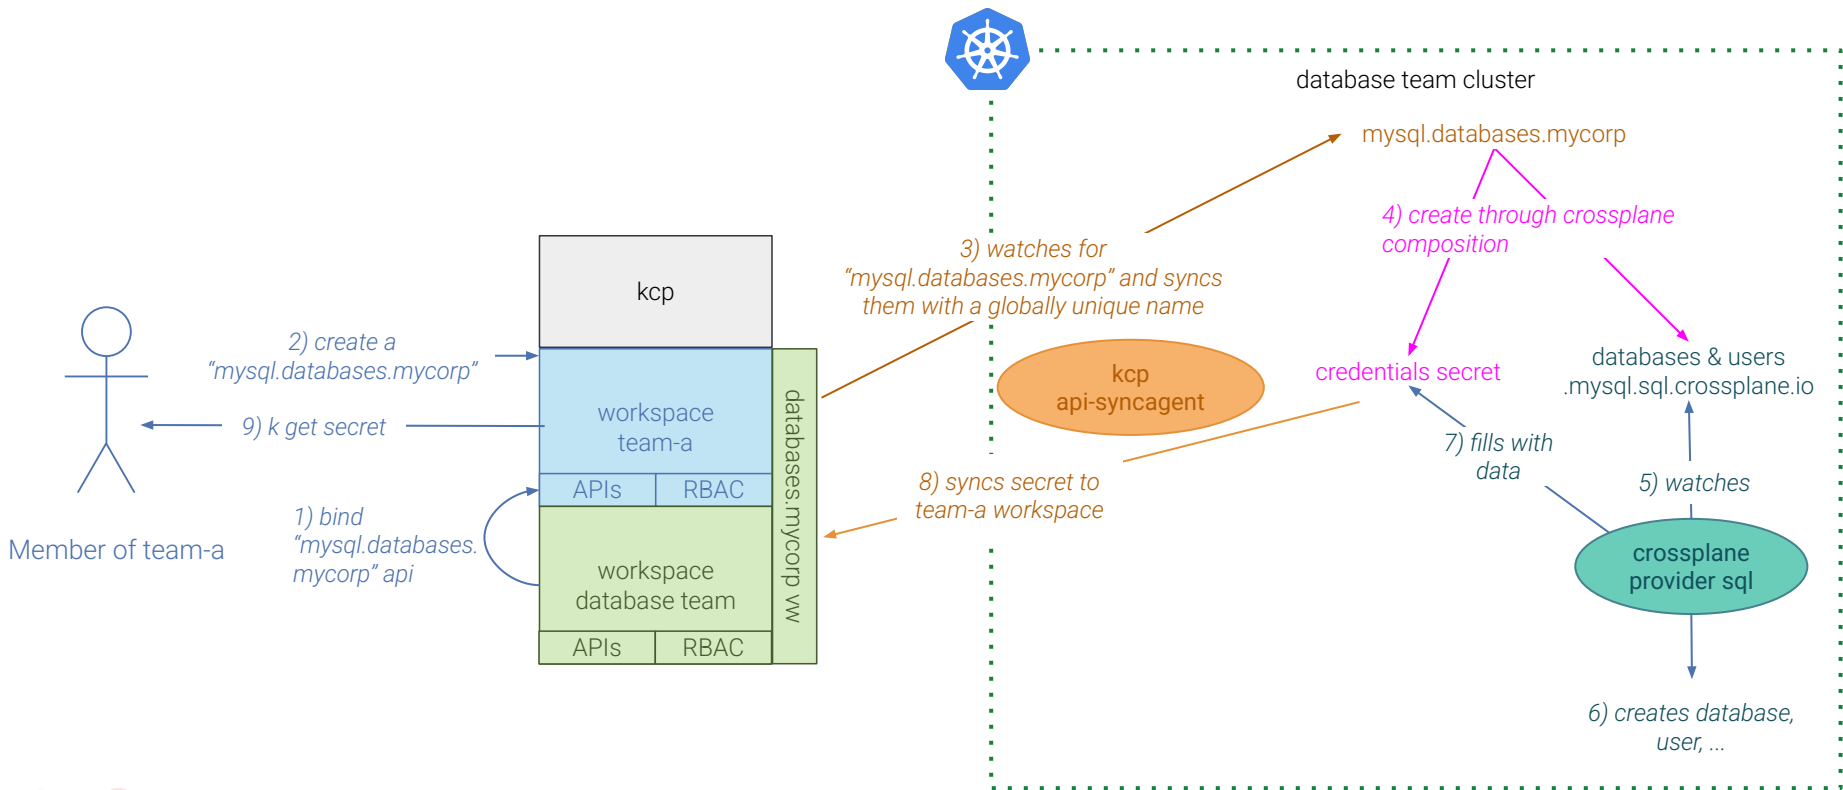
# How can we achieve this?



+



# Putting the Pieces Together



# Enough Slides, Let's Hop In



# Why so complicated?

## Multi Tenancy Degrees of Freedom

kcp			
workspace team-a		workspace team-b	
CRDs	RBAC	CRDs	RBAC
workspace database team		workspace certificate team	
CRDs	RBAC	CRDs	RBAC



database team cluster

One **or** multiple  
provider instances



certificate team cluster

One **or** multiple  
provider instances





# Multi Tenancy Degrees of Freedom

- Multiple consumer teams can have their own bound APIs via kcp workspaces
- Multiple providers can be consumed per workspace
- Multiple consumers can be served by a single crossplane provider instance



# One last thing about kcp



# So you have decided you want to try kcp out

```
$ kubectl ws .
```

```
Current workspace is 'root'.
```

```
$ kubectl get crd
```

```
No resources found
```

```
$ kubectl api-resources
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
workspaces	ws	tenancy.kcp.io/v1alpha1	false	Workspace
workspacetypes		tenancy.kcp.io/v1alpha1	false	WorkspaceType
databases		databases.demo.example.com	false	Database
...				

```
$ kubectl explain workspaces
```

```
# will return workspaces api definition
```



# Wrapping Up



# Wrapping Up



By using Crossplane v2 and kcp we can create true Multi Tenancy with multiple degrees of freedom

<https://crossplane.io>

[github.com/lsviben](https://github.com/lsviben)

[linkedin.com/in/lovr-o-sviben](https://linkedin.com/in/lovr-o-sviben)

**kcp and crossplane are community projects! We welcome everyone to build the future together.**



<https://kcp.io>

[kcp-users & kcp-dev on Kubernetes Slack](#)

[github.com/SimonTheLeg](https://github.com/SimonTheLeg)

[linkedin.com/in/simon-bein](https://linkedin.com/in/simon-bein)




# Backup: RBAC Extension for APIBindings

Binding to exported APIs  
requires RBAC permissions  
on the **APIExport**.

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: bind-apiexport
rules:
- apiGroups:
  - apis.kcp.io
  resources:
  - apiexports
  verbs:
  - use
  resourceNames:
  - demo.embik.me
```



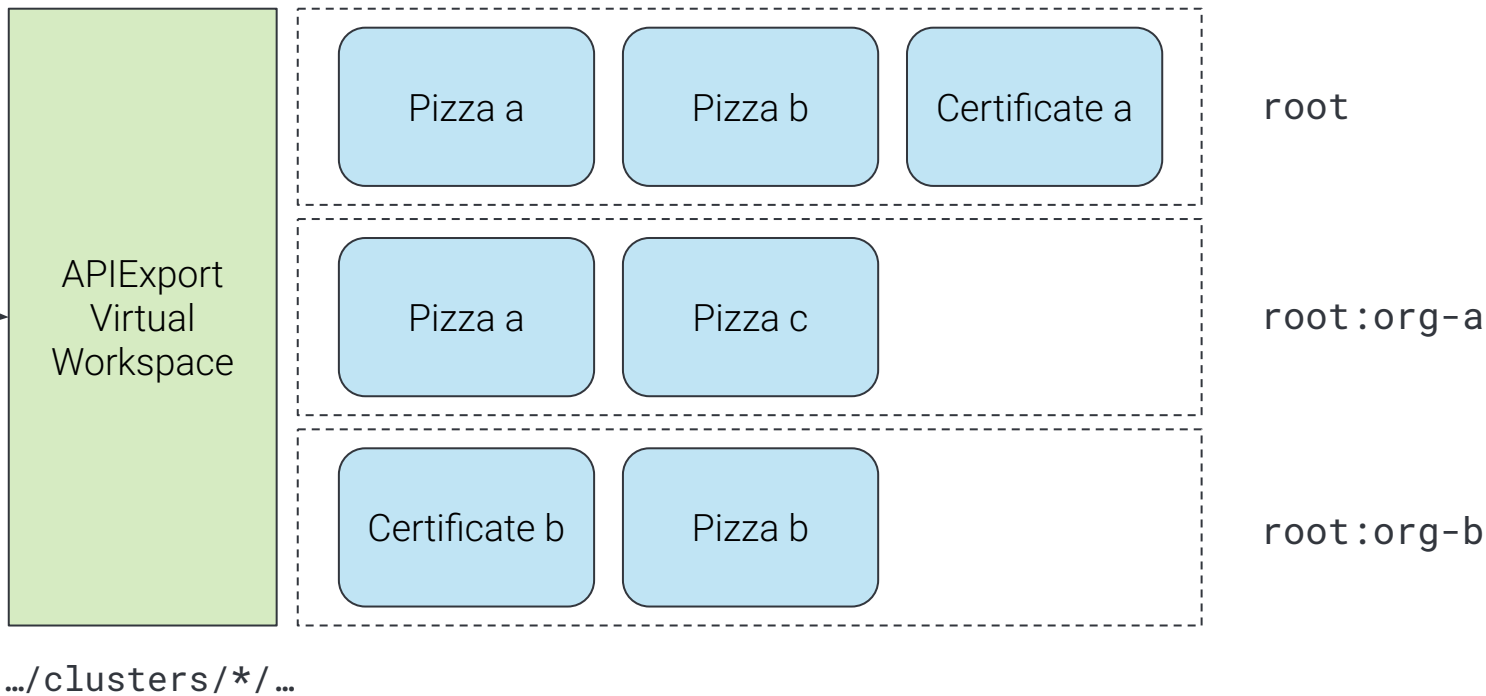
# Backup: Virtual Workspaces for Controllers



Proxy that provides a computed view



Access requires special RBAC!





# Backup: How to Build a KCP-aware Controller

## 1 Use kcp-aware client and cache

```
MapperProvider: kcp.NewClusterAwareMapperProvider,  
NewClient:      kcp.NewClusterAwareClient,  
NewCache:       kcp.NewClusterAwareCache,  
NewAPIReader:   kcp.NewClusterAwareAPIReader,
```

## 2

## Reconcile in Virtual Workspace via **Cluster**

```
sigs.k8s.io/controller-runtime/pkg/cluster.Cluster
```

## 3 Reconcile with logical cluster in context

```
ctx = kontext.WithCluster(ctx, logicalcluster.Name(request.ClusterName))
```